
Convolutional Neural Networks using separable filters

Petrescu Viviana
EPFL
viviana.petrescu@epfl.ch

Abstract

CNN are the state-of-the-art machine learning techniques which achieved best results on various computer vision tasks ranging from the large scale ImageNet object recognition challenge to segmentation in bio medical imaging. This technical report presents initial results on using CNN with separable filters for speeding up the testing time.

1 Introduction

Although proven to be very powerful, CNN are much slower for both training and testing than their counter parts SVM or Random Forests. In the forward pass, the computational complexity of evaluating one image of size $W \times H$ with J filters of size $d_1 \times d_2$ is $O(WHJd_1d_2)$.

Although improving the training time would be very beneficial since it would allow for more configurations to be tried and for larger networks, increasing effort has been put also into speeding up only testing time, since training can be done offline.

2 Speeding up CNN

By far the most common approach for speeding up CNNs is to run them on the GPU, making use of the parallelism nature of the algorithm. Alternatively or combined, FFT can be used for the convolutional operations (for both training and testing)[3]. For testing, a significant speedup can be obtained in hardware by using FGPAs [2]. The major drawback of FGPAs is that they are harder to program for people who do not work in the field.

In [1] they prove speedup of CNNs for two different schemes using separable filters, one of which is similar with our approach but using a different optimization algorithm for obtaining the separable filters. If the 2d filters are decomposed into a set of separable 1d filters of rank K , the complexity per image becomes $O(WH(J + d_1 + d_2))$. Thus, we obtain a speedup if $K \ll \frac{Jd_1d_2}{J+d_1+d_2}$.

3 Separable filters

In our approach, the set of filters χ for one convolutional layer is approximated as a set of separable filters, as shown in Fig1.

We can obtain an approximation by minimizing the equation:

$$\underset{a,b,c}{\text{minimize}} \quad \left\| \chi - \sum_{r=1}^{r=R} a_r \circ b_r \circ c_r \right\|_2^2$$

We do this using a Matlab implementation of the the Canonical Polyadic decomposition (a generalization of SVD for tensors), algorithm which uses non linear conjugate gradient method. The




Figure 1: Tensor decomposition

| Tensor size | Typical rank |
|---|--------------|
| $I \times J \times K$ with $JK \leq I$ (very tall) | JK |
| $I \times J \times K$ with $JK - J < I < JK$ (tall) | I |
| $I \times J \times K$ with $I = JK - J$ (compact) | $I, I + 1$ |

Table 1: Typical rank for certain tensor shapes

optimization framework requires as input the rank R , which is almost never known. When the rank used for decomposition is not the theoretical one, the decomposition is not very accurate and the algorithm needs to be run with different initialization parameter until it converges. What is known is an upper bound on the theoretical rank R for a general tensor 3d tensor $\chi \in \mathbb{R}^{I \times J \times K}$, given by:

$$\text{rank}(\chi) \leq \min\{IJ, JK, KI\}$$

Besides the theoretical rank, there is also the notion of typical rank, which is any rank that appears with probability greater than 0, or that it is most common [cite]. Some of the known typical ranks are shown in Table 1. Looking at the first row of Table 1, we conclude that for a 'very tall' set of filters like the ones present in large CNNs, we should expect the rank to be equal with the product of the two kernel dimensions.

4 Experiments

We run the experiments using Theano library in Python. We varied the rank for every convolutional layer and compared the separable version with the non separable one. We recorded the time, the change in performance and how well the separable filters approximate the 2d filters. We did this for the MNIST and for the Mitochondria Striatum datasets.

4.1 MNIST

MNIST consists of a curated set of black and white images (28x28) depicting handwritten digits. The dataset is split into training 60.000 samples, validation 10.000 and testing 10.000 samples. For this set, the approach of [?] using CNNs is the best winning model with a 0.23 error rate (equivalent with 23 out of 10.000 digits not recognized correctly). We start our experience with the reference Theano model for MNIST which achieves 0.82 error rate. Its configuration is shown in Fig2

4.1.1 CNN Model 1

We first decompose the first convolutional layer1. From the theoretical point of view, according to Table 1, we have a compact tensor $R^{20 \times 5 \times 5}$ with a typical rank of 20 or 21, so we expect a very good approximation for rank in that range. Fig[ref] shows how well is the approximation with varying rank from 4 to 16 in steps of 2 on the x axis and the fit on the y axis (100 corresponding to perfect fit). As expected, the fit is almost perfect for high ranks and decreases afterwards.

Using separable filters, we obtain a theoretical speedup for conv layer 1 if $K \ll \frac{Jd_1d_2}{J+d_1+d_2} = 16.66$. Fig3a shows the time improvements using separable filters for conv layer 1. The blue line represents

| Layer | Type | Maps and neurons | Kernel size |
|-------|------------------|------------------|-------------|
| 0 | input | 1 map of 28x28 | |
| 1 | convolutional | 20 maps of 24x24 | 5x5 |
| 2 | max pooling | 20 maps of 12x12 | |
| 3 | convolutional | 50 maps of 8x8 | 5x5 |
| 4 | max pooling | 50 maps of 4x4 | |
| 5 | fully conntected | 500 | |
| 6 | fully conntected | 2 neurons | |

Table 2: Small CNN for MNIST set

the time per layer using non separable filters. In our implementation, the separable filters give a speedup for any rank, as the red curve is below the blue one. We note that that theoretically the non separable time (blue line) and the separable conv with rank 16 should have given the same running time. The reason for the difference is due to implementation details and language specific.




Figure 2: Rank versus Fit

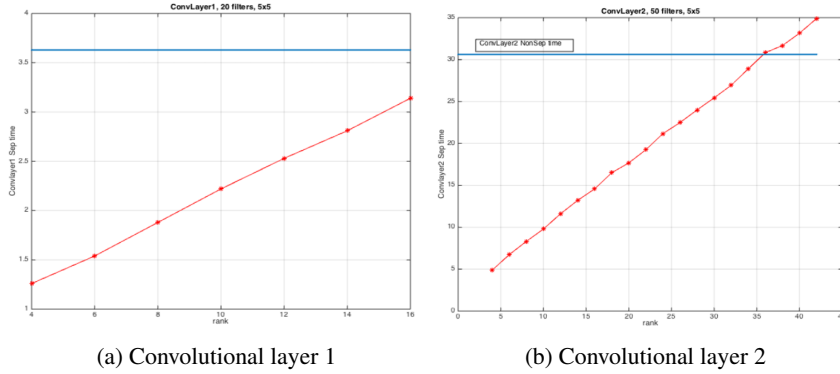


Figure 3: Rank vs Time

The question that remains now is to see how far can we afford to reduce the rank such that we do not lose much from the accuracy. Fig [cite]a shows how the performance of the CNN drops with decreasing rank. With rank 16 we do not lose at all performance, while for rank between 10 and 16 the error rate stays between 0.8 and 0.9 per cent, which is quite low. According to the application, even a rank of 8 or 6 is acceptable, since the drop is not more than 1 per cent.

We then kept the first layer fixed, and approximated the second layer. Similarly as before, according to Table 1, we have a very tall tensor $R^{50 \times 5 \times 5}$ with a typical rank of 25. Fig2b shows how well is the approximation with varying rank from 4 to 42 in steps of 2. As expected, the fit is almost perfect (99.9 fit) for rank higher than the theoretical one of 25.




Figure 4: Convolutional layer 2

| Layer | Type | Maps and neurons | Kernel size |
|-------|-----------------|------------------|-------------|
| 0 | input | 1 map of 28x28 | |
| 1 | convolutional | 20 maps of 24x24 | 9x9 |
| 2 | max pooling | 20 maps of 12x12 | |
| 3 | convolutional | 50 maps of 8x8 | 9x9 |
| 4 | max pooling | 50 maps of 4x4 | |
| 5 | fully connected | 500 | |
| 6 | fully connected | 2 neurons | |

Table 3: Bigger CNN for MNIST set

From the complexity perspective, the separable cnn is faster if we use a rank $K \ll \frac{Jd_1d_2}{J+d_1+d_2} = 20.83$. In our case, for rank = 20 we obtained almost a 40% speedup. This is due as before due to language specific implementation. For rank 10, using separable filters is actually 3 times faster (theoretically 2 times faster). If we keep the rank greater than 10, the recognition never drops more than 1 error rate. For rank 12 the fit is 79 but the error rate is still almost unchanged which means we can easily use rank with not so perfect fit and still obtain good performance. The first layer with separable filters of similar fit of 79 obtain a slightly lower performance.

4.1.2 CNN Model 2

We use next a model with filters of size 9×9 for both layers.

The theoretical rank is less than 49 and we obtain speedup if the rank $K \ll 59.5$ The results can be seen in the figure below




Figure 5: Convolutional layer 2

| Layer | Type | Maps and neurons | Kernel size |
|-------|------------------|------------------|-------------|
| 0 | input | 1 map of 51x51 | |
| 1 | convolutional | 10 maps of 46x46 | 6x6 |
| 2 | max pooling | 10 maps of 23x23 | |
| 3 | convolutional | 20 maps of 18x18 | 6x6 |
| 4 | max pooling | 20 maps of 9x9 | |
| 3 | convolutional | 50 maps of 4x4 | 6x6 |
| 4 | max pooling | 50 maps of 2x2 | |
| 5 | fully conntected | 100 | |
| 6 | fully conntected | 2 neurons | |

Table 4: CNN for Mitochondria set

4.2 Mitochondria

For the EM Mitochondria set, CNN are employed to solve the segmentation problem. Every pixel is classified as being part of the mitochondria or not. For pixel classification, a patch of 51x51 surrounding the pixel is extracted from the image and fed as input to a CNN which classifies it. For determining the size and parameters of the CNN I used a set of 100.000 samples for training and 20.000 for validation (containing half positive and half negative samples). After determining the best setup, the network was trained on a larger set of 1million samples for training and 200.000 for validation.

The final test data represented a 3D volume consisting in 318 slices of 400x661 images. Thus, one frame 264.400 datasamples and all frames 97 million datasamples.

4.3 Pre trained ImageNet CNN

Here we look to see if we could gain something from CNNs of large networks. The network below is the reference net from Alex Krizhevski paper (without the split to be run on two GPUs). The net was trained by [cite] here.

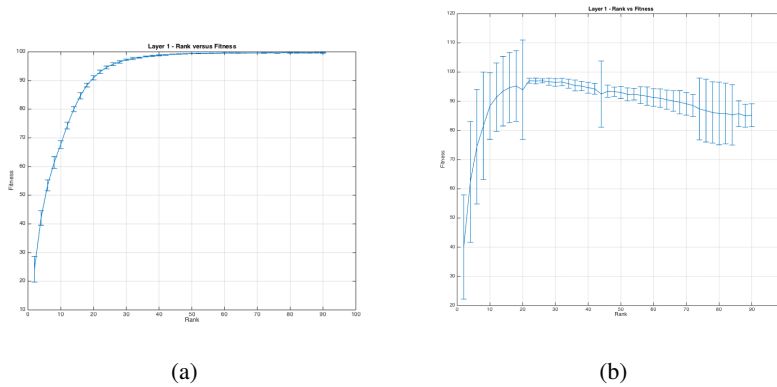


Figure 6: ImageNet Layers

5 Conclusions

Theoretical bounds of the separable filters choice are proven. Improved recognition on mitochondria set.




Figure 7: ImageNet Layers

Acknowledgments

We would like to thank Simon Arosi that helped me a lot during this project giving me constant feedback.

References

- [1] M. Jaderberg, A. Vedaldi, and A. Zisserman. Speeding up convolutional neural networks with low rank expansions. In *British Machine Vision Conference*, 2014.
- [2] Yann LeCun, Koray Kavukcuoglu, and Clément Farabet. Convolutional networks and applications in vision. In *Circuits and Systems (ISCAS), Proceedings of 2010 IEEE International Symposium on*, pages 253–256. IEEE, 2010.
- [3] Michaël Mathieu, Mikael Henaff, and Yann LeCun. Fast training of convolutional networks through ffts. *CoRR*, abs/1312.5851, 2013.